

CS 35L Software Construction Lab

Week 4 – Debugging

Debugging Process

- Reproduce the bug
- Simplify program input
- Use a debugger to track down the origin of the problem
- Fix the problem

Debugger

- A program that is used to run and debug other (target) programs
- Advantages:
 - Programmer can:
 - step through source code line by line
 - each line is executed on demand
 - interact with and inspect program at run-time
 - If program crashes, the debugger outputs where and why it crashed

GDB – GNU Debugger

- Debugger for several languages
 - C, C++, Java, Objective-C... more
- Allows you to inspect what the program is doing at a certain point during execution
- Logical errors and segmentation faults are easier to find with the help of gdb

Using GDB

1. Compile Program

- Normally: `$ gcc [flags] <source files> -o <output file>`
- Debugging: `$ gcc [other flags] -g <source files> -o <output file>`
 - enables built-in debugging support

2. Specify Program to Debug

- `$ gdb <executable>`
- `$ gdb`
- `(gdb) file <executable>`

Run-Time Errors

- Segmentation fault
 - Program received signal SIGSEGV, Segmentation fault.
0x0000000000400524 in *function* (arr=0x7fffc902a270, r1=2, c1=5, r2=4, c2=6) at *file.c*:12
 - Line number where it crashed and parameters to the function that caused the error
- Logic Error
 - Program will run and exit successfully
- How do we find bugs?

Using GDB

3. Run Program

- `(gdb) run` or
- `(gdb) run [arguments]`

4. In GDB Interactive Shell

- Tab to Autocomplete, up-down arrows to recall history
- `help [command]` to get more info about a command

5. Exit the gdb Debugger

- `(gdb) quit`

Setting Breakpoints

- Breakpoints
 - used to stop the running program at a specific point
 - If the program reaches that location when running, it will pause and prompt you for another command
- Example:
 - (gdb) `break file1.c:6`
 - Program will pause when it reaches line 6 of file1.c
 - (gdb) `break my_function`
 - Program will pause at the first line of `my_function` every time it is called
 - (gdb) `break [position] if expression`
 - Program will pause at specified position only when the expression evaluates to true

Breakpoints

- Setting a breakpoint and running the program will stop program where you tell it to
- You can set as many breakpoints as you want
 - (gdb) `info breakpoints | break | br | b`
shows a list of all breakpoints

Basic commands

- `(gdb) step` - Step to next line of code. Will step into a function.
- `(gdb) print <var>` - Print value stored in variable.
- `(gdb) next` - Execute next line of code. Will not enter functions.
- `(gdb) continue` - Continue execution to next break point.
- `(gdb) set var <name>=<value>` - Executes rest of program with new value of variable.

Deleting, Disabling and Ignoring BPs

- (gdb) delete [bp_number | range]
 - Deletes the specified breakpoint or range of breakpoints
- (gdb) disable [*bp_number* | *range*]
 - Temporarily deactivates a breakpoint or a range of breakpoints
- (gdb) enable [*bp_number* | *range*]
 - Restores disabled breakpoints
- If no arguments are provided to the above commands, all breakpoints are affected!!
- (gdb) ignore *bp_number iterations*
 - Instructs GDB to pass over a breakpoint without stopping a certain number of times.
 - bp_number: the number of a breakpoint
 - Iterations: the number of times you want it to be passed over

Displaying Data

- Why would we want to interrupt execution?
 - to see data of interest at run-time:
 - (gdb) `print [/format] expression`
 - Prints the value of the specified expression in the specified format
 - Formats:
 - d: Decimal notation (default format for integers)
 - x: Hexadecimal notation
 - o: Octal notation
 - t: Binary notation

Resuming Execution After a Break

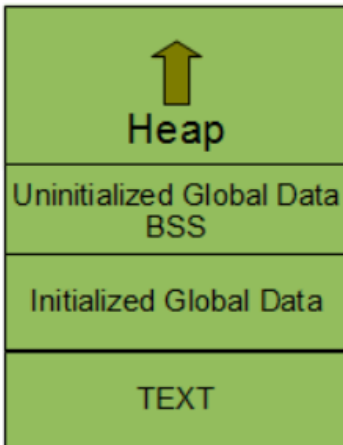

- When a program stops at a breakpoint
 - 4 possible kinds of gdb operations:
 - **c or continue**: debugger will continue executing until next breakpoint
 - **s or step**: debugger will continue to next source line
 - **n or next**: debugger will continue to next source line in the current (innermost) stack frame
 - **f or finish**: debugger will resume execution until the current function returns. Execution stops immediately after the program flow returns to the function's caller
 - the function's return value and the line containing the next statement are displayed

Watchpoints

- Watch/observe changes to variables
 - (gdb) `watch my_var`
 - sets a watchpoint on `my_var`
 - the debugger will stop the program when the value of *my_var* changes
 - old and new values will be printed
 - (gdb) `rwatch expression`
 - The debugger stops the program whenever the program reads the value of any object involved in the evaluation of *expression*

Process Memory Layout

(Higher Address)



(Lower Address)

- TEXT segment
 - Contains machine instructions to be executed
- Global Variables
 - Initialized
 - Uninitialized
- Heap segment
 - Dynamic memory allocation
 - malloc, free
- Stack segment
 - Push frame: Function invoked
 - Pop frame: Function returned
 - Stores
 - Local variables
 - Return address, registers, etc
- Command Line arguments and Environment Variables

Stack Info

- A program is made up of one or more functions which interact by calling each other
- Every time a function is called, an area of memory is set aside for it. This area of memory is called a **stack frame** and holds the following crucial info:
 - storage space for all the local variables
 - the memory address to return to when the called function returns
 - the arguments, or parameters, of the called function
- Each function call gets its own stack frame. Collectively, all the stack frames make up the **call stack**

Stack Frames and the Stack

```
1  #include <stdio.h>
2  void first_function(void);
3  void second_function(int);
4
5  int main(void)
6  {
7      printf("hello world\n");
8      first_function();
9      printf("goodbye goodbye\n");
10
11     return 0;
12 }
13
14
15 void first_function(void)
16 {
17     int imidate = 3;
18     char broiled = 'c';
19     void *where_prohibited = NULL;
20
21     second_function(imidate);
22     imidate = 10;
23 }
24
25
26 void second_function(int a)
27 {
28     int b = a;
29 }
```

Frame for `main()`

Frame for `first_function()`

Return to `main()`, line 9

Storage space for an int

Storage space for a char

Storage space for a void *

Frame for `second_function()`:

Return to `first_function()`, line 22

Storage space for an int

Storage for the int parameter named `a`

Analyzing the Stack in GDB

- `(gdb) backtrace | bt`
 - Shows the call trace (the call stack)
 - Without function calls:
 - #0 main () at program.c:10
 - one frame on the stack, numbered 0, and it belongs to main()
 - After call to function `display()`
 - #0 display (z=5, zptr=0xbffffb34) at program.c:15
 - #1 0x08048455 in main () at program.c:10
 - Two stack frames: frame 1 belonging to main() and frame 0 belonging to display().
 - Each frame listing gives
 - the arguments to that function
 - the line number that's currently being executed within that frame

Analyzing the Stack

- (gdb) info frame
 - Displays information about the current stack frame, including its return address and saved register values
- (gdb) info locals
 - Lists the local variables of the function corresponding to the stack frame, with their current values
- (gdb) info args
 - List the argument values of the corresponding function call

Other Useful Commands

- (gdb) info functions
 - Lists all functions in the program
- (gdb) list
 - Lists source code lines around the current line

Lab 4

- Download old version of coreutils with buggy ls program
 - Untar, configure, make
- Bug: ls -lt mishandles files whose time stamps are very far in the past. It seems to act as if they are in the future

```
$ touch -d '1918-11-11 11:00 GMT' wwi-armistice
$ touch now
$ sleep 1
$ touch now1
$ ls -lt wwi-armistice now now1
```

Output:

```
-rw-r--r-- 1 eggert eggert 0 Nov 11 1918 wwi-armistice
-rw-r--r-- 1 eggert eggert 0 Feb 5 15:57 now1
-rw-r--r-- 1 eggert eggert 0 Feb 5 15:57 now
```

Goal: Fix the Bug

- **Reproduce the Bug**
 - Follow steps on lab web page
- **Simplify input**
 - Run ls with -l and -t options only
- **Debug**
 - Use gdb to figure out what's wrong
 - \$ gdb ./ls
 - (gdb) run -lt wwi-armistice now now1
(run from the directory where the compiled ls lives)
- **Patch**
 - Construct a patch "lab5.diff" containing your fix
 - It should contain a ChangeLog entry followed by the output of diff -u

Lab Hints

- Don't forget to answer all questions! (lab4.txt)
- Make sure not to submit a reverse patch! (lab4.diff)
- “Try to reproduce the problem in your home directory, instead of the \$tmp directory. How well does SEASnet do?”
 - Timestamps represented as seconds since Unix Epoch 1970.1
 - SEASnet NFS filesystem has unsigned 32-bit time stamps
 - Local File System on Linux server has signed 32-bit time stamps
 - If you touch the files on the NFS filesystem it will return timestamp around 2054
 - => files have to be touched on local filesystem (df -l)
- Use “info functions” to look for relevant starting point
- Compiler optimizations: -O2 -> -O0
 - ./configure CFLAGS="...-O0"