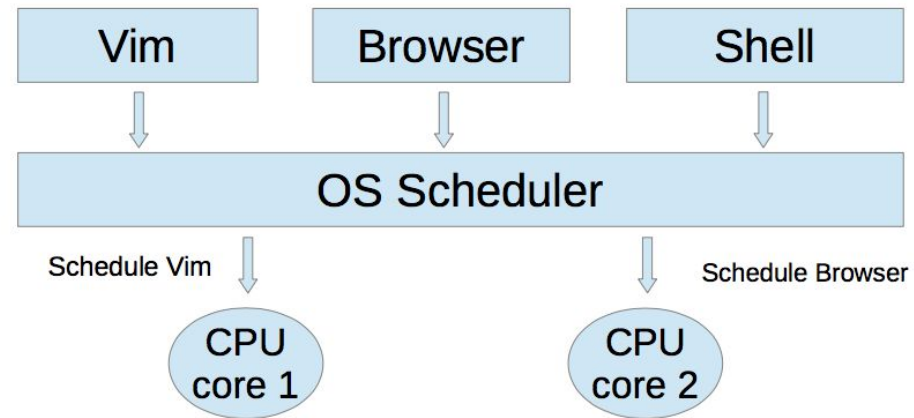
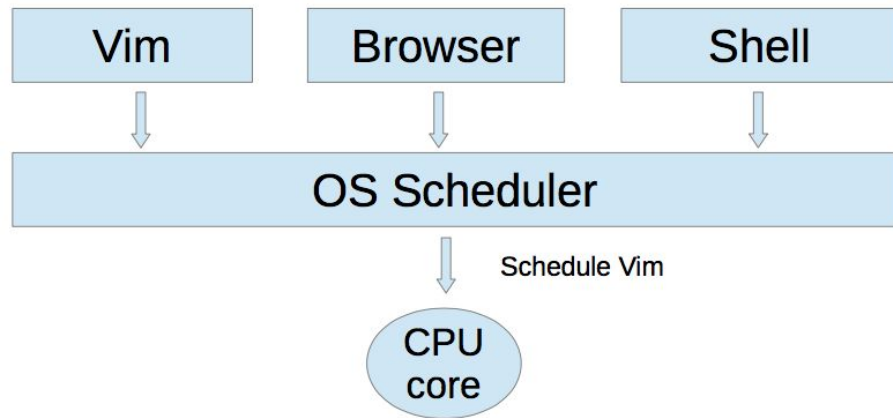


Multithreading/Parallel Processing

Week 6

Multitasking

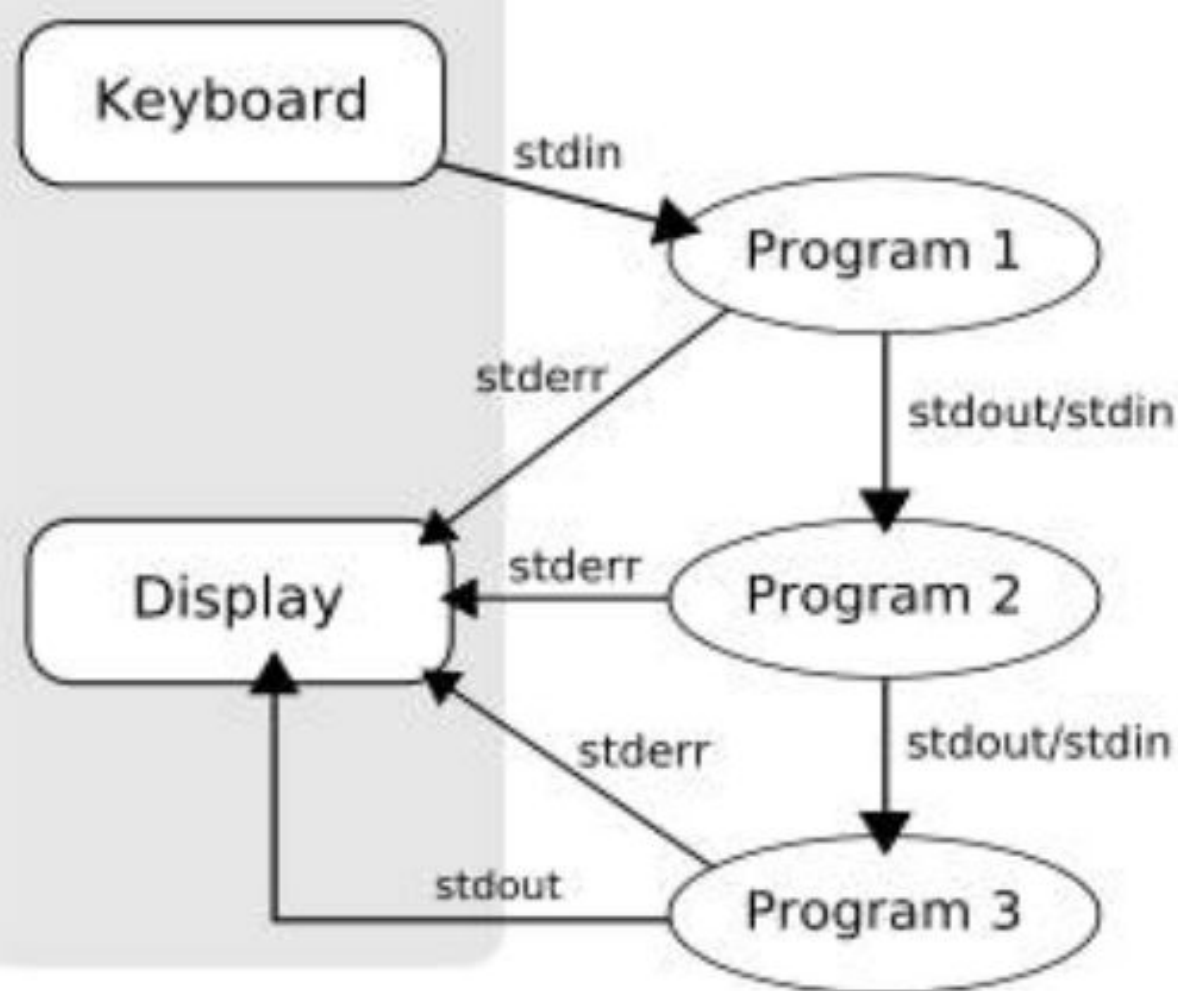
- Run multiple processes **simultaneously** to increase performance
- Processes do not share internal structures (stacks, globals, etc)
 - Communicate via **IPC** (inter-process communication) methods
 - Pipes, sockets, signals, message queues
- **Single core: Illusion** of parallelism by switching processes quickly (**time-sharing**). [Why is illusion good?](#)
- **Multi-core: True** parallelism. Multiple processes execute **concurrently** on different CPU cores



Multitasking


- `tr -s '[:space:]' '\n' | sort -u | comm -23 - words`
- Three separate processes spawned simultaneously
 - P1 - `tr`
 - P2 - `sort`
 - P3 - `comm`
- Common buffers (pipes) exist between 2 processes for communication
 - '`tr`' writes its stdout to a buffer that is read by '`sort`'
 - '`sort`' can execute, as and when data is available in the buffer
 - Similarly, a buffer is used for communicating between '`sort`' and '`comm`'

Text terminal

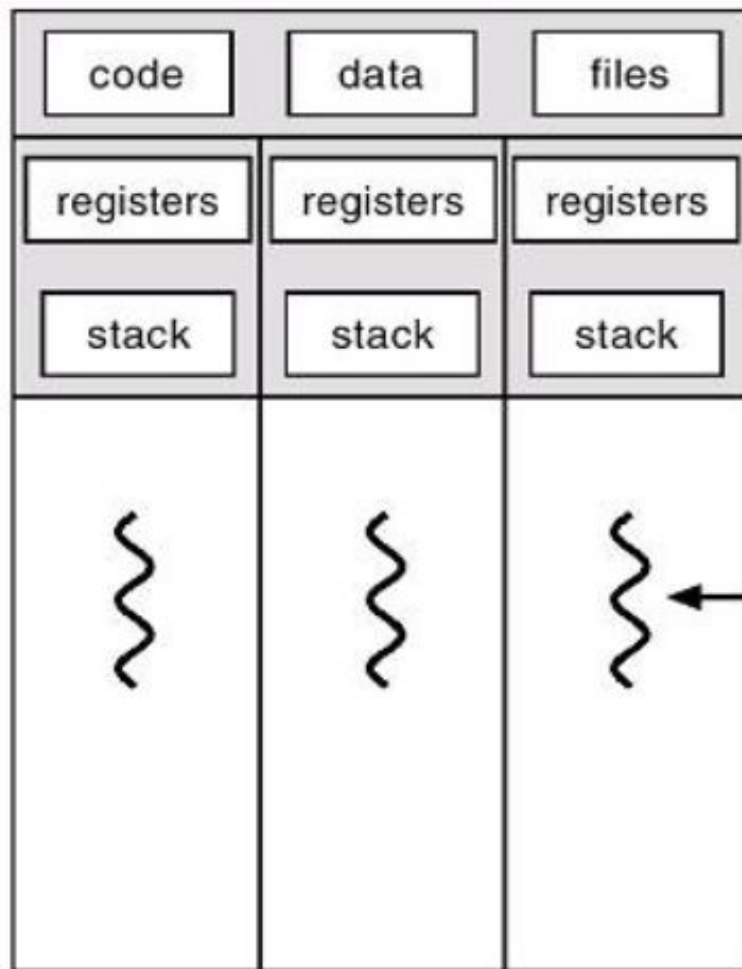


Threads

- A flow of instructions, path of execution within a process
- It is a basic unit of CPU utilization
- Each thread has its own:
 - Stack
 - Registers
 - Thread ID
- Each thread shares the following with other threads belonging to the same process
 - Code
 - Heap
 - Global Data
 - OS resources (files,I/O)
- A process can be single-threaded or multi-threaded
- Threads in a process can run in parallel
(provide another type of parallelism)



single-threaded




multithreaded

Single threaded execution


```
int global_counter = 0  
int main()  
{  
    ...  
    foo(arg1,arg2);  
    bar(arg3,arg4,arg5);  
    ...  
    return 0;  
}
```

```
void foo(arg1,arg2)  
{  
    //code for foo  
}  
void bar(arg3,arg4,arg5)  
{  
    //code for bar  
}
```



Sequential execution of subroutines

Multiple threads sharing a single CPU



Multiple threads on multiple CPUs



Multi threaded execution (single core)

```
int global_counter = 0
```

```
int main()
```

```
{
```

```
...
```

```
foo(arg1,arg2);
```

```
bar(arg3,arg4,arg5);
```

```
...
```

```
return 0;
```

```
}
```

```
void foo(arg1,arg2)
```

```
{
```

```
//code for foo
```

```
}
```

```
void bar(arg3,arg4,arg5)
```

```
{
```

```
//code for bar
```

```
}
```



Time Sharing – Illusion of multithreaded parallelism
(Thread switching has less overhead compared to process switching)

Multi threaded execution (multiple cores)

```
int global_counter = 0
```

```
int main()
```

```
{
```

```
...
```

```
foo(arg1,arg2);
```

```
bar(arg3,arg4,arg5);
```

```
...
```

```
return 0;
```

```
}
```

```
void foo(arg1,arg2)
```

```
{
```

```
//code for foo
```


```
}
```

```
void bar(arg3,arg4,arg5)
```

```
{
```

```
//code for bar
```

```
}
```



True multithreaded parallelism

Multithreading properties


- Efficient way to **parallelize** tasks
- **Thread switches are less expensive** compared to process switches (context switching)
- Inter-thread communication is easy, via **shared global** data (heap)
- Need **synchronization** among threads accessing same data

Shared Memory

- Makes multithreaded programming
 - **Powerful**
can easily access data and share it among threads
 - **More efficient**
No need for system calls when sharing data
Thread creation and destruction less expensive than process creation and destruction
 - **Non-trivial**
Have to prevent several threads from accessing and changing the same shared data at the same time (synchronization)

Race Condition

```
int count = 0;  
void increment()  
{  
    count = count + 1;  
}
```



Result depends on order of execution
=> Synchronization needed

Lab 6

- Evaluate the performance of multithreaded 'sort' command
 - **od -An -f -N 4000000 < /dev/urandom | tr -s ' ' '\n' > random.txt**
 - Might have to modify the command above
- Delete the empty line
 - **time -p sort -g --parallel=2 numbers.txt > /dev/null**
- Add /usr/local/cs/bin to PATH
 - **\$ export PATH=/usr/local/cs/bin:\$PATH**
- Generate a file containing 10M random **single-precision floating point numbers**, one per line with no white space
 - /dev/urandom: pseudo-random number generator

Lab 6

- od
 - write the contents of its input files to standard output in a user-specified format
 - Options
 - -t f: Double-precision floating point
 - -N <count>: Format no more than *count* bytes of input
- sed, tr
 - Remove address, delete spaces, add newlines between each float

Lab 6

- use `time -p` to time the command `sort -g` on the data you generated
- Send output to `/dev/null`
- Run `sort` with the `--parallel` option and the `-g` option: compare by general numeric value
 - Use `time` command to record the real, user and system time when running `sort` with 1, 2, 4, and 8 threads
 - `$ time -p sort -g file_name > /dev/null` (1 thread)
 - `$ time -p sort -g --parallel=[2, 4, or 8] file_name > /dev/null`
 - Record the times and steps in `log.txt`